

NBpy: Network-based (R)Statistics in Python

Ljuba, Waleed, & Zeus

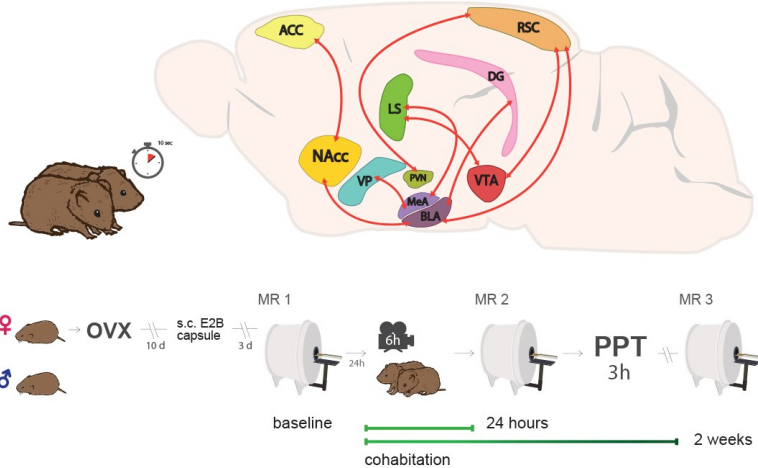
July 30, 2021

The logo for NeuroHackAcademy is displayed on a dark gray background with a hexagonal pattern. The text "NEUROHACKADEMY" is written in a bold, white, sans-serif font.

NEUROHACKADEMY

Context

Longitudinal brain networks



(López-Gutierrez et al., *eLife*, 2021)

Context

Prairie voles

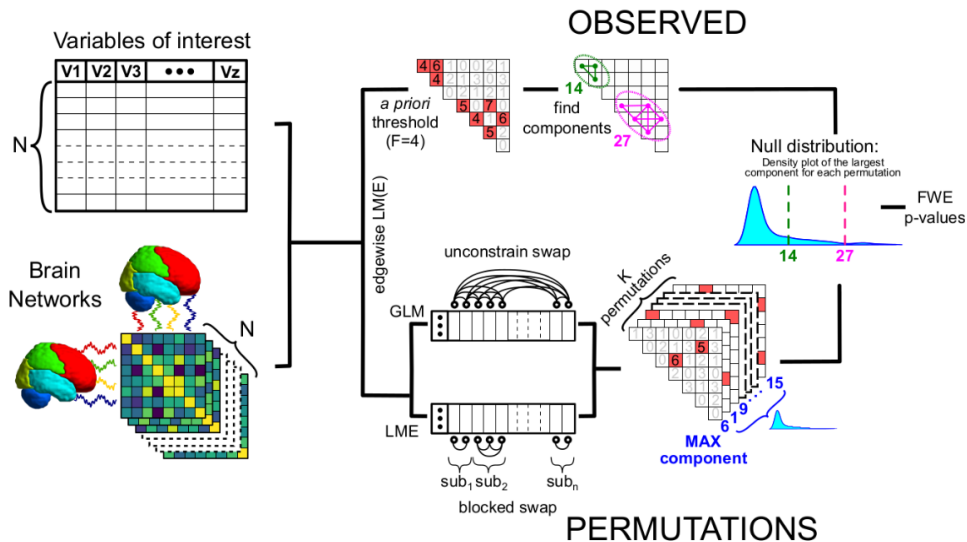
NBpy

Intro

NBR



Network Based Statistics framework



(Gracia-Tabuenca & Alcauter, *bioRxiv*, 2021)

- ▶ First CRAN release 0.1.2 (March 2020)
 - ▶ <https://cran.r-project.org/package=NBR>
 - ▶ >6k downloads
(<https://cranlogs.r-pkg.org/badges/grand-total/NBR>)
 - ▶ Computing efficiency is the biggest challenge, according to the mails from various users.

Python implementation

- ▶ **statsmodels** is a python framework designed for statistics, and allows users to fit statistical models using R-style formulas.
- ▶ Simply import statsmodels' R-API:

```
import statsmodels.formula.api as smf
```

- ▶ And start R-hacking in Python, with API customization.

statsmodel implementation

- Statsmodel execution for 1000 randomly selected endog columns finished on average: 3.22 ± 0.73 [ms]

```
def fit_linear_model_stats(data, endog, exog_relation):  
    """  
    Generate linear model with statsmodels.ols.  
    """  
  
    lm = smf.ols(  
        formula='%s ~ %s' % (endog, exog_relation),  
        data=data  
    ).fit()  
  
    df_result = pd.concat([  
        lm.pvalues,  
        lm.tvalues,  
    ], axis=1)  
    df_result.columns = ['pvalues_%s' % endog, 'tvalues_%s' % endog]  
    return df_result.drop(index=['Intercept'])
```


- Statsmodel execution for 1000 randomly selected endog columns finished on average: 3.22 ± 0.73 [ms]

```
def fit_linear_model_stats(data, endog, exog_relation):  
    """  
    Generate linear model with statsmodels.ols.  
    """  
  
    lm = smf.ols(  
        formula='%s ~ %s' % (endog, exog_relation),  
        data=data  
    ).fit()  
  
    df_result = pd.concat([  
        lm.pvalues,  
        lm.tvalues,  
    ], axis=1)  
    df_result.columns = ['pvalues_%s' % endog, 'tvalues_%s' % endog]  
    return df_result.drop(index=['Intercept'])
```

- If we run linear regression across all endog columns (interconnections)

```
duration = []
for endog in random_endogs[:20]:
    st = time.time()
    dfa = []
    for col in data.drop(columns=predictor_cols).columns:
        df = fit_linear_model_stats(data, col, 'Group + Sex*Age')
        dfa.append(df)
    dfa = pd.concat(dfa, axis=1)
    et = time.time()
    dur = et - st
    duration.append(dur)
print('Statsmodel execution for all ndog columns finished on average: %.2f ± %.2f [s]')
```

Statsmodel execution for all ndog columns finished on average: 1.24 ± 0.02 [s]

- ▶ Can we improve performance by computing linear models across all input endog columns, by utilizing matrix computations?
- ▶ For that reason we will call Pythonic API to handle input data with statsmodels.OLS.

```
class statsmodels.regression.linear_model.OLS(endog, exog=None, missing='none', hasconst=None, **kwargs)
```

[\[source\]](#)

Ordinary Least Squares

Parameters

endog : [array_like](#)

A 1-d endogenous response variable. The dependent variable.

exog : [array_like](#)

A nobs x k array where nobs is the number of observations and k is the number of regressors. An intercept is not included by default and should be added by the user. See `statsmodels.tools.add_constant`.

- ▶ Not possible to have an endog input with dimension higher than 1.

ValueError: shapes (48,378) and (48,378) not aligned: 378 (dim 1) != 48 (dim 0)

sklearn/numpy

- ▶ sklearn/numpy to the rescue.
- ▶ We can perform matrix/broadcasting operations on a lowest level of abstraction.

```
def fit_linear_model(X, y):
    X = data_preprocessing(data)
    y = data.drop(columns=predictor_cols)
    # Step 1: train model
    lr_model = linear_model.LinearRegression().fit(X,y)
    # Step 2: prediction
    predictions = lr_model.predict(X)
    # Step 3: pval and tval calculation
    # 3a: get coefficients and stack them with intercept
    lr_parameters = np.vstack((lr_model.intercept_.T, lr_model.coef_.T)).T
    # 3b: append 1s to input data (for intercept)
    X_intercept = np.append(np.ones((len(X), 1)), X, axis=1)
    # 3b: get mean squared error between true values and predictions and scale it
    mse = np.sum((y - predictions) ** 2, axis=0) / (X_intercept.shape[0] - X_intercept.shape[1])
    # 3c: calculate invariant of input dataset (this might be tricky to do for larger matrix)
    X_inv = np.linalg.inv(np.dot(X_intercept.T, X_intercept)).diagonal()
    # 3d: estimate variance and standard deviation to calculate tvals
    var = np.dot(mse.values.reshape(mse.shape[0], 1), X_inv.reshape(1, X_inv.shape[0]))
    std = np.sqrt(var)
    tvals = lr_parameters / std
    # 3e: calculate pvalues
    pvals = [2 * (1 - scipy.stats.t.cdf(np.abs(i), (X_intercept.shape[0] - X_intercept.shape[1]))) for i in tvals]
    # 3f: save it in dataframe
    index = ['Intercept'] + X.columns.tolist()
    columns = ['tvals_%s' % c for c in y.columns] + ['pvals_%s' % c for c in y.columns]
    df_result = pd.DataFrame(
        np.vstack((tvals, pvals)).T,
        columns=columns,
        index=index,
    ).iloc[1:]

    return df_result
```

- ▶ If we run linear regression across all endog columns (interconnections)
- ▶ We achieve improvement of calculating linear regression models for all input endog columns from 1.24 seconds to 56.6ms

```
: %%time  
df_res = fit_linear_model(X,y)
```

CPU times: user 70.9 ms, sys: 127 ms, total: 198 ms

Wall time: 56.6 ms

sklearn/numpy

- ▶ Although statsmodel has better support for statistical calculations (ie. faster computation of p and t vals), p and t vals calculation is embedeed into package.

Feature	sklearn	statsmodels
P/T vals embedeed	NO	YES
R alike formula	NO	YES
Preprocessing categorical columns	YES	NO
Speed [ms]	3.3	4.1
Broadcasting	YES	NO

Results comparison

\$GroupPatient

Component	strn	strnFWE
5	0.260	0.5101727
6	0.439	0.2882506
7	0.034	5.4706903

\$SexM

Component	strn	strnFWE
13	0.45	0.002298164
14	0.45	0.027913618

\$Age

Component	strn	strnFWE
5	0.407	0.1326686

\$`SexM:Age`

NULL

\$GroupPatient

Component	strn	strnFWE
4	0.221	0.5101727
5	0.309	0.2882506
6	0.004	5.4706903

\$SexM

Component	strn	strnFWE
12	0.495	0.002298164
13	0.476	0.027913618

\$Age

Component	strn	strnFWE
4	0.397	0.1326686

\$`SexM:Age`

NULL

Conclusion

- ▶ R implementation has reach statistical libraries.
- ▶ On contrary, Python requires implementation of p and t values.
- ▶ However, it is faster to load bigger datasets with Python and it allows broadcasting/matrix operations.
- ▶ **Improvements:**
 - ▶ Instead of **sklearn** we can utilize **numpy.linalg.lstsq** library to compute coefficients of linear regression.
 - ▶ Improve calculation of **FWE** strength values.